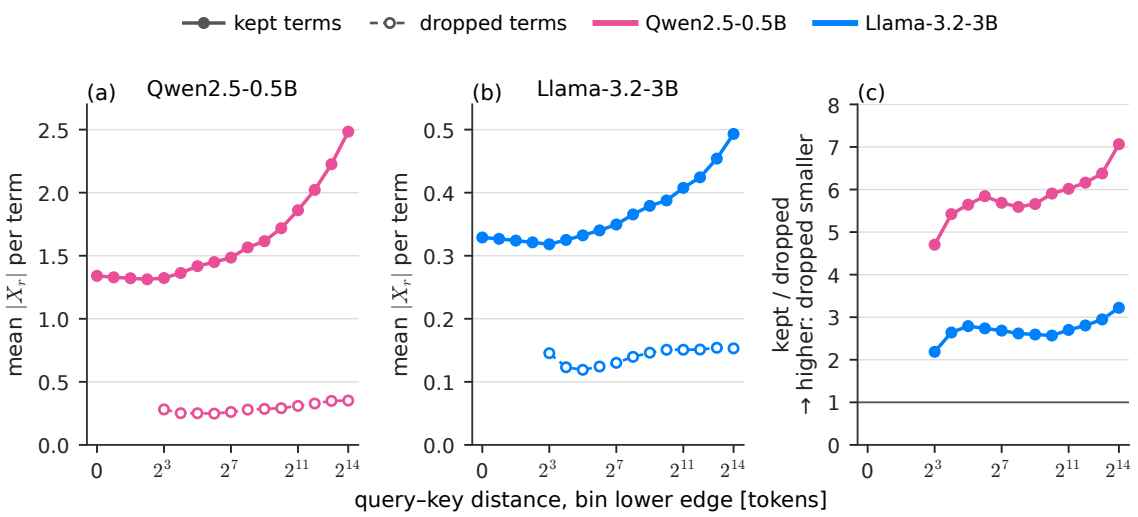
kept terms
dropped terms
Qwen2.5-0.5B
Llama-3.2-3B
(a) Qwen2.5-0.5B
(b) Llama-3.2-3B
(c)
mean $|X_r|$ per term
mean $|X_r|$ per term
kept / dropped
→ higher: dropped smaller
query–key distance, bin lower edge [tokens]
0
2^3
2^7
2^{11}
2^{14}