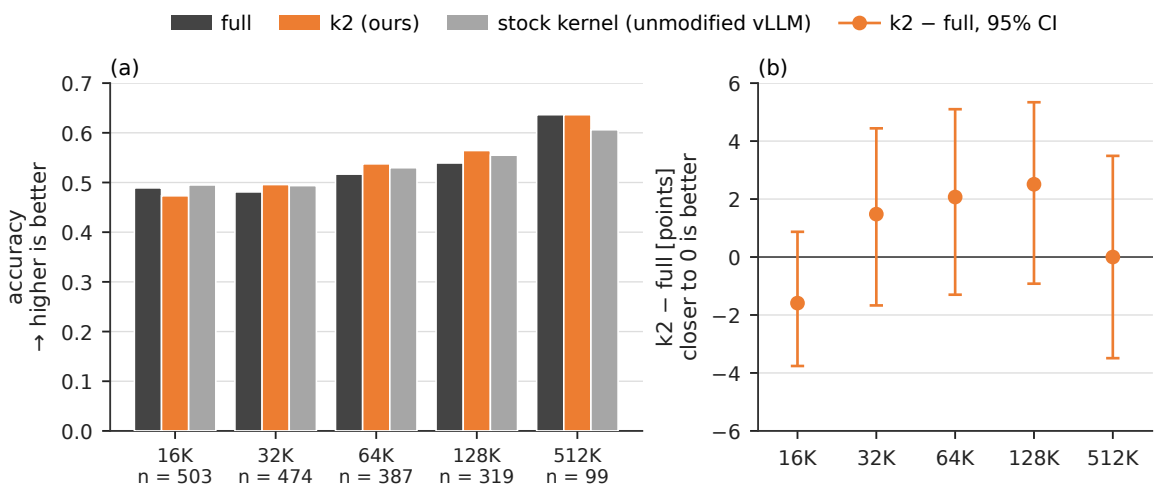
full
k2 (ours)
stock kernel (unmodified vLLM)
k2 − full, 95% CI
(a)
accuracy
→ higher is better
0.7
0.6
0.5
0.4
0.3
0.2
0.1
0.0
16K
n = 503
32K
n = 474
64K
n = 387
128K
n = 319
512K
n = 99
(b)
k2 − full [points]
closer to 0 is better
6
4
2
0
−2
−4
−6
16K
32K
64K
128K
512K