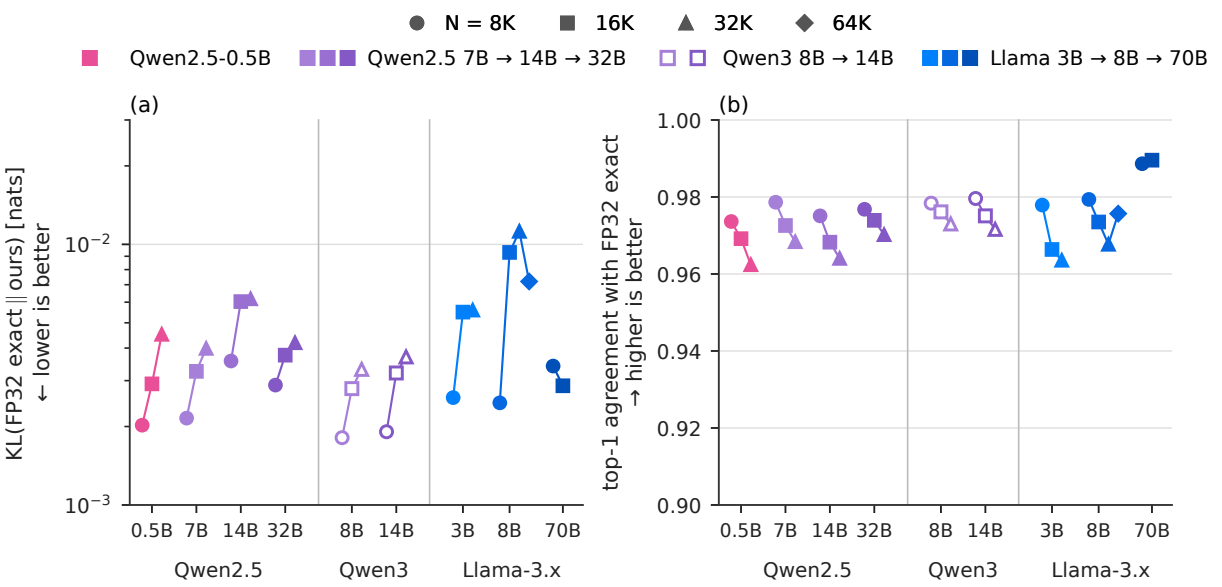
N = 8K
16K
32K
64K
Qwen2.5-0.5B
Qwen2.5 7B → 14B → 32B
Qwen3 8B → 14B
Llama 3B → 8B → 70B
(a)
KL(FP32 exact ‖ ours) [nats]
← lower is better
10^{-2}
10^{-3}
0.5B 7B 14B 32B 8B 14B 3B 8B 70B
Qwen2.5 Qwen3 Llama-3.x
(b)
top-1 agreement with FP32 exact
→ higher is better
1.00
0.98
0.96
0.94
0.92
0.90
0.5B 7B 14B 32B 8B 14B 3B 8B 70B
Qwen2.5 Qwen3 Llama-3.x