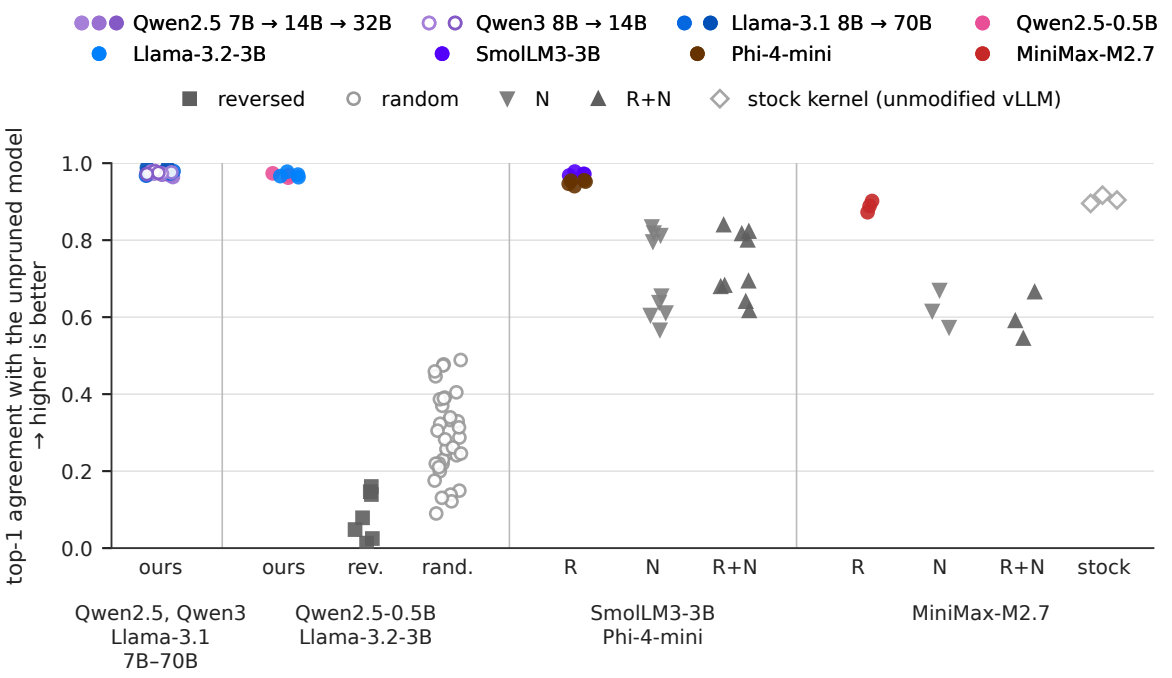
Qwen2.5 7B → 14B → 32B
Qwen3 8B → 14B
Llama-3.1 8B → 70B
Qwen2.5-0.5B
Llama-3.2-3B
SmolLM3-3B
Phi-4-mini
MiniMax-M2.7
reversed
random
N
R+N
stock kernel (unmodified vLLM)
top-1 agreement with the unpruned model
→ higher is better
1.0
0.8
0.6
0.4
0.2
0.0
ours
ours
rev.
rand.
R
N
R+N
R
N
R+N
stock
Qwen2.5, Qwen3
Llama-3.1
7B–70B
Qwen2.5-0.5B
Llama-3.2-3B
SmolLM3-3B
Phi-4-mini
MiniMax-M2.7